
Operator-Adapted Reproducing Kernel Geometries for Structure-Preserving Representation Learning

Himanshu Singh^{✉*}

Department of Meteorology and Atmospheric Science
College of Earth and Mineral Sciences
The Pennsylvania State University
University Park, PA 16802
hfs5641@psu.edu

Abstract

The representation geometry determines which transformations appear to preserve, distort, or separate the represented objects. In most representation-learning settings, this geometry is fixed in advance or learned from data, while the action of an operator is subsequently studied within it. We consider the inverse problem by asking, given a linear operator, can one construct a reproducing kernel Hilbert space (RKHS) geometry in which its action is exactly norm preserving? We establish necessary and sufficient conditions for the existence of such an operator-adapted RKHS in finite-dimensional function spaces and characterize the complete family of reproducing kernels realizing this invariance. The resulting classification shows that existence is governed by the spectral structure of the operator, while the spectral multiplicities determine the degrees of freedom available in the invariant representation geometry. In particular, an operator that is not unitary under the original geometry may become unitary after an appropriate change of RKHS geometry, whereas operators with incompatible spectral structure admit no such positive-definite realization. These results provide a mathematical framework for constructing representations from the prescribed operator structure and suggest an alternative principle for structure-preserving representation learning where rather than modifying an operator to respect a predetermined geometry, one may construct the geometry to respect the operator.

1 Introduction

The performance of any Artificial Intelligence and machine learning model is heavily dependent on the choice of data representation [2] that is technically often referred to as *features* [2]. In this concern, representation learning plays a monumental role by providing a priori details from the aspect of several general purposes such as probability distributions embedding [7], spatial and temporal coherence, manifold, semi-supervised learning, sparsity, etc.

Kernel machine learning [5, 18] offers a promising path for representation learning of features for high dimensional data. Essentially, the high dimensional data gets encoded in the inner-product between the infinite dimensional nonlinear kernel feature space over the underlying Hilbert space. These specialized function spaces corresponding to kernels are called *reproducing kernel Hilbert spaces* [1] (RKHS). Besides being able to handle the huge dimension of data, RKHS also offers a unique notion of representation geometry that also encapsulates data similarity, distance and structural variation by getting embedded in it through the kernel representation.

*<https://himanshuvnm.github.io/>

Linear transformations such as the Koopman (*composition*) operators [12], are unitary in L^2 [17] under measure-preserving invertible dynamics [13], making them desirable operator learning for achieving robust prediction, control, regression and etc. in various scientific domains [21, 4, 9]. In general, it is immediate to observe that the same *desirable* notion cannot be offered within the RKHS.

In this paper, we are *motivated* by this very aim of adapting RKHS for which the operator norm geometry is already known. In particular, we establish an existence and classification theory for such operator-adapted RKHS geometries. Exact preservation is possible precisely when the matrix representation of the operator is diagonalizable with spectrum on the unit circle. When this condition is satisfied, the complete family of admissible geometries is determined by the spectral decomposition of the operator. We further show how the same principle produces an operator-adapted Gaussian RKHS by replacing the Euclidean quadratic form in the Gaussian radial basis function kernel [5, 20] with an invariant positive-definite metric. This construction permits transformations that are not unitary in Euclidean geometry to induce unitary Koopman composition operators in the adapted RKHS whenever the corresponding invariant metric exists.

Problem statement. Let X be a nonempty subset inside $\mathbb{R}^d$, and let $\mathbb{C}^X := \{f \mid f : X \rightarrow \mathbb{C}\}$. Let $1 \leq j \leq m$ be natural indexes and $\psi_j : X \rightarrow \mathbb{C}$ be a basis function with $\mathcal{V} := \text{span}\{\psi_1, \dots, \psi_m\} \subseteq \mathbb{C}^X$. Let $T : \mathcal{V} \rightarrow \mathcal{V}$ be a linear operator. We are interested in inner products on $\mathcal{V}$ for which $(\mathcal{V}, \langle \cdot, \cdot \rangle_{\mathcal{H}_k})$ is an RKHS $\mathcal{H}_k$ with a reproducing kernel function k such that T acts isometrically:

$$\|Tf\|_{\mathcal{H}_k} = \|f\|_{\mathcal{H}_k} \quad \forall f \in \mathcal{V}.$$

Accordingly, we address three questions:- (i) under what conditions does such an RKHS geometry exist; (ii) when it exists, how can the complete family of admissible kernels be characterized; and (iii) how can the resulting invariant geometry be realized within a standard kernel family so that a prescribed transformation induces a unitary composition action?

1.1 Related Work

Transformation-invariant kernels [14, 18] have been studied for geometric actions such as translations, rotations, and dilations, while group-averaging constructions have been used to obtain invariant feature maps and associated invariant RKHSs [19, 16]. In these approaches, the transformation class is typically specified first, and invariance is built into the kernel or feature representation.

Recent results characterize when mappings of the input space induce bounded composition operators in RKHSs associated with analytic positive definite functions; in broad classes, boundedness can itself force the underlying map to be affine [10, 8]. Such results typically begin with a fixed RKHS and ask which composition operators [6] it admits. Our finite-dimensional problem reverses this direction, that is, the operator is fixed, while the Hilbertian geometry is unknown and is to be constructed so that the operator becomes isometric.

Kernel methods have also become important in Koopman operator theory, where the RKHS structure is used to approximate Koopman operators, eigenfunctions, or observables and to derive approximation guarantees [21, 11]. These works generally assume or select an RKHS and then study the Koopman operator within that space. Our construction instead asks whether the geometry itself can be chosen from the operator.

2 Operator-Adapted Reproducing Kernel Geometries

2.1 RKHS

Let $k : X \times X \rightarrow \mathbb{C}$ be a continuous Hermitian positive definite kernel denoted by $\succ 0$. At $x \in X$, define the canonical feature $\phi_z(x) = k(x, z)$ centered at $z \in X$. Then, the linear space given by $\text{span}\{\phi_z : z \in X\}$ extends by completion to a Hilbert space $\mathcal{H}_k$, called the *reproducing kernel Hilbert space* (RKHS) [5] generated by $k(\cdot, \cdot)$. For all functions $f \in \mathcal{H}_k$, the following properties hold true at $x \in X$:

$$\langle f, k(\cdot, x) \rangle_{\mathcal{H}_k} = f(x) \tag{1a}$$

$$\langle k(\cdot, y), k(\cdot, x) \rangle_{\mathcal{H}_k} = k(x, y) \tag{1b}$$

$$|f(x)| \leq \|f\|_{\mathcal{H}_k} \cdot (k(x, x))^{1/2}. \tag{1c}$$

The property shown in (1a) is called the *reproducing property* of $k(\cdot, x)$ and therefore $k(\cdot, x)$ is called the reproducing kernel of $\mathcal{H}_k$. When $f = k(\cdot, y)$ in (1a), we get (1b) and finally (1c) shows the boundedness of the *evaluation functional* [15, 1] $\mathcal{E}_x(f) := f(x)$ for $f \in \mathcal{H}_k$ in $x \in X$.

2.2 Operator Adapted RKHS

Recall that X is a nonempty set and let $\mathcal{V} = \text{span}\{\psi_1, \dots, \psi_m\}$ be a m -dimensional space of complex-valued functions on X , where $\psi_1, \dots, \psi_m$ are linearly independent such that $\dim \mathcal{V} = m$. Define the vector valued representation as follows through linearly independent basis functions ψ_j :

$$\Psi(x)^* = [\psi_1(x) \quad \dots \quad \psi_m(x)], \quad x \in X,$$

where $\bullet^*$ implies the Hermitian-adjoint. For $c := [c_j]_{j=1}^m \in \mathbb{C}^m$, write $f_c(x) = \Psi(x)^* c = \sum_{j=1}^m c_j \psi_j(x)$. Let $T : \mathcal{V} \rightarrow \mathcal{V}$ be a linear operator. Relative to the prescribed basis, let $M \in \mathbb{C}^{m \times m}$ denote its matrix representation, so that $Tf_c = f_{Mc}$.

We write $\mathbb{H}_{++}^m := \{G \in \mathbb{C}^{m \times m} : G = G^*, G \succ 0\}$ for the open convex cone of Hermitian positive-definite matrices [3]. For $G \in \mathbb{H}_{++}^m$, we defined the induced inner-product and norm as follows:

$$\begin{aligned} \langle f_c, f_d \rangle_G &= d^* G c \\ \|f_c\|_G^2 &= c^* G c. \end{aligned} \tag{2}$$

In this section, we characterize those RKHS geometries on $\mathcal{V}$ for which T acts isometrically. The following proposition provides the RKHS realization whose inner product is given in (2).

Proposition 2.1. *For every $G \in \mathbb{H}_{++}^m$, the Hilbert space $(\mathcal{V}, \langle \cdot, \cdot \rangle_G)$ is an RKHS denoted by $\mathcal{H}_{k_G}$ with the reproducing kernel given as: $k_G(x, y) = \Psi(x)^* G^{-1} \Psi(y)$, $x, y \in X$.*

We also have uniqueness of kernel parametrization, meaning if $k_{G_1}(x, y) = k_{G_2}(x, y)$ for all $x, y \in X$, then $G_1 = G_2$. Refer Proposition B.2 for this result. Now that we have established the preliminary foundation structure of RKHS $\mathcal{H}_{k_G}$ in Proposition 2.1, we then establish the existence of an invariant RKHS geometry for a linear operator T as follows.

Theorem 2.2. *Let $T : \mathcal{V} \rightarrow \mathcal{V}$ and $G \in \mathbb{H}_{++}^m$. Consider RKHS $\mathcal{H}_{k_G}$ as given in Proposition 2.1, then the following are equivalent:*

1. T acts isometrically on $\mathcal{H}_{k_G}$;
2. $\|Tf\|_{\mathcal{H}_{k_G}} = \|f\|_{\mathcal{H}_{k_G}} \quad \forall f \in \mathcal{V}$;
3. $M^* G M = G$.

Consequently, $\exists G \in \mathbb{H}_{++}^m$ for which T acts isometrically if and only if M is similar to a unitary matrix. Equivalently, M is diagonalizable and $\sigma(M) \subset \mathbb{T}$, where $\mathbb{T} := \{z \in \mathbb{C} : |z| = 1\}$.

Remark 2.3. *Condition $\sigma(M) \subset \mathbb{T}$ alone is not sufficient. A nontrivial Jordan block associated with a unimodular eigenvalue cannot satisfy $M^* G M = G$ for any $G \succ 0$. Thus, diagonalizability constitutes an additional obstruction to the existence of an exact invariant positive-definite geometry.*

For $M \in \mathbb{C}^{m \times m}$, we let the collection of all admissible matrices $G \in \mathbb{H}_{++}^m$ such that $M^* G M = G$ holds; this is given by $\mathcal{G}(M) := \{G \in \mathbb{H}_{++}^m : M^* G M = G\}$. The following theorem helps to provide the classification of invariant geometries whose equivalent conditions were provided in Theorem 2.2. Define the family of operator-adapted kernels by $\mathcal{R}_\Psi(T) := \{k_G : G \in \mathcal{G}(M)\}$.

Theorem 2.4. *Assume that M is diagonalizable and its spectrum satisfies $\sigma(M) \subset \mathbb{T}$. Let $\mu_1, \dots, \mu_r$ denote the distinct eigenvalues of M , with multiplicities $m_1, \dots, m_r$, respectively. Choose an eigenvector matrix S such that $M = SDS^{-1}$, where $D = \text{diag}(\mu_1 I_{m_1}, \dots, \mu_r I_{m_r})$. Then $G \in \mathbb{H}_{++}^m$ belongs to $\mathcal{G}(M)$ if and only if $G = S^{-*} \text{diag}(H_1, \dots, H_r) S^{-1}$, where $H_j \in \mathbb{H}_{++}^{m_j}$, $j = 1, \dots, r$. Here, $S^{-*} := (S^{-1})^* = (S^*)^{-1}$.*

We take advantage of the spectral structure of $M \in \mathbb{C}^{m \times m}$ to classify exactly which RKHS geometries and, consequently, which RKHS kernels are compatible with operator invariance; this classification is completed in the following corollary.

Corollary 2.5. According to the hypotheses of [Theorem 2.4](#), the complete family of reproducing kernels for which T acts isometrically is $k_{\mathbf{H}}(x, y) = \Psi(x)^* S \operatorname{diag}(H_1^{-1}, \dots, H_r^{-1}) S^* \Psi(y)$, where $(H_1, \dots, H_r) \in \prod_{j=1}^r \mathbb{H}_{++}^{m_j}$.

Corollary 2.6. According to the hypotheses of [Theorem 2.4](#), $\dim_{\mathbb{R}} \mathcal{G}(M) = \dim_{\mathbb{R}} \mathfrak{K}_{\Psi}(T) = \sum_{j=1}^r m_j^2$. Thus, the spectral multiplicity determines the degrees of freedom in the family of operator-adapted RKHS kernels.

2.3 Numerical Verification of Operator-Adapted Geometry

We illustrate the construction on a $d = 20$ linear representation operator M that is similar to a unitary matrix but is not unitary in the native Euclidean geometry. Thus, although its eigenvalues lie on $\mathbb{T}$, its singular values differ substantially from one, showing that spectral preservation alone does not imply preservation of the ambient representation geometry. We construct $G \succ 0$ satisfying $M^* G M = G$. Equivalently, the similarity-transformed operator $\tilde{M} = G^{1/2} M G^{-1/2}$ is unitary.

[Figure 1](#) illustrates the resulting change in geometry. The spectrum is unchanged under the similarity transformation ([Figure 1a](#)), while the singular values of M collapse to one despite the substantial singular value variation of M in the Euclidean metric ([Figure 1b](#)). Then, we compare the Gaussian radial basis function kernel $k_0(x, y) = \exp\left(-\frac{\|x-y\|_2^2}{2\gamma^2}\right)$ with its operator-adapted counterpart $k_G(x, y) = \exp\left(-\frac{(x-y)^* G (x-y)}{2\gamma^2}\right)$ with the kernel bandwidth $\gamma = 3$.

To test invariance directly, we sample representations $z_1, \dots, z_N \in \mathbb{R}^{20}$, with $N = 120$ and compare pairwise kernel values before and after application of M . For a kernel k , we define the entrywise invariance discrepancy as $\Delta_k(i, j) = |k(Mz_i, Mz_j) - k(z_i, z_j)|$. Thus, $\Delta_k(i, j) = 0$ for every pair (i, j) whenever k is exactly invariant under M . We summarize the discrepancy by the normalized error given as follows:

$$\varepsilon_k := \sqrt{\frac{\sum_{i,j=1}^N \Delta_k^2(i, j)}{\sum_{i,j=1}^N |k(z_i, z_j)|^2}} = \sqrt{\frac{\sum_{i,j=1}^N |k(Mz_i, Mz_j) - k(z_i, z_j)|^2}{\sum_{i,j=1}^N |k(z_i, z_j)|^2}}.$$




Figure 1: Operator adaptation converts a spectrally stable but geometrically non-isometric representation into an *exact* isometry.

Using a common color scale, [Figure 1c-d](#) demonstrates the collapse of the invariance error from $\varepsilon_{k_0} = 5.61 \times 10^{-1}$ to $\varepsilon_{k_G} = 1.12 \times 10^{-15}$, while the operator spectrum remains unchanged. Thus, adapting the RKHS metric does not alter the operator spectrum but changes the geometry in which the same operator acts from non-isometric to isometric.

3 Conclusion

We formulated operator-adapted RKHS construction as an inverse geometric problem and characterized exactly when a finite-dimensional linear operator admits an invariant positive-definite geometry. Existence is equivalent to similarity to a unitary matrix, while spectral multiplicities parameterize the complete family of invariant RKHS kernels. The same principle yields operator-adapted Gaussian kernels that preserve pairwise similarity even when the underlying transformation is non-unitary.

in Euclidean geometry. Thus, rather than modifying an operator to fit a prescribed representation geometry, the geometry itself can be constructed to respect the operator.

As a limitation of our proposed research work, the present results concern finite-dimensional function spaces. The corresponding problem in infinite-dimensional RKHSs leads to the operator equation $T^*GT = G$, with G a bounded, positive, and boundedly invertible operator, together with the question of how such invariant Hilbert geometries are realized at the level of reproducing kernels. Establishing an intrinsic classification in this setting is a natural direction beyond the finite-dimensional theory developed here.

References

- [1] Nachman Aronszajn. Theory of reproducing kernels. *Transactions of the American Mathematical Society*, 68(3):337–404, 1950. [1](#), [3](#)
- [2] Yoshua Bengio, Aaron Courville, and Pascal Vincent. Representation Learning: A Review and New Perspectives. *IEEE Trans. Pattern Anal. Mach. Intell.*, 35(8):1798–1828, August 2013. [1](#)
- [3] Rajendra Bhatia. *Positive Definite Matrices*. Princeton University Press, 2009. [3](#)
- [4] Lea Bold, Friedrich M Philipp, Manuel Schaller, and Karl Worthmann. Kernel-based Koopman approximants for control: Flexible sampling, error analysis, and stability. *SIAM Journal on Control and Optimization*, 63(6):4044–4071, 2025. [2](#)
- [5] Andreas Christmann and Ingo Steinwart. Support Vector Machines. *Springer*, 2008. [1](#), [2](#), [9](#)
- [6] Jr Cowen and Barbara I MacCluer. *Composition operators on spaces of analytic functions*. Routledge, 2019. [2](#)
- [7] Kenji Fukumizu, Arthur Gretton, Gert Lanckriet, Bernhard Schölkopf, and Bharath K Sriperumbudur. Kernel choice and classifiability for RKHS embeddings of probability distributions. *Advances in Neural Information Processing Systems*, 22, 2009. [1](#)
- [8] Efrain Gonzalez, Moad Abudia, Michael Jury, Rushikesh Kamalapurkar, and Joel A Rosenfeld. The Kernel Perspective on Dynamic Mode Decomposition. *Transactions on Machine Learning Research*, 2024. [2](#)
- [9] ADG Hales, MJ Colbrook, and C Jiang. Improving dynamic mode decomposition of tandem cylinder flow with nonlinear dictionaries. *Physics of Fluids*, 37(6), 2025. [2](#)
- [10] Masahiro Ikeda, Isao Ishikawa, and Yoshihiro Sawano. Boundedness of composition operators on reproducing kernel Hilbert spaces with analytic positive definite functions. *Journal of Mathematical Analysis and Applications*, 511(1):126048, 2022. [2](#)
- [11] Yoshinobu Kawahara. Dynamic mode decomposition with reproducing kernels for Koopman spectral analysis. *Advances in Neural Information Processing Systems*, 29, 2016. [2](#)
- [12] Bernard O Koopman. Hamiltonian systems and transformation in Hilbert space. *Proceedings of the National Academy of Sciences*, 17(5):315–318, 1931. [2](#)
- [13] Andrzej Lasota and Michael C Mackey. *Chaos, fractals, and noise: stochastic aspects of dynamics*. Springer Science & Business Media, 2013. [2](#)
- [14] Fabien Lauer and Gérard Bloch. Incorporating prior knowledge in support vector machines for classification: A review. *Neurocomputing*, 71(7-9):1578–1594, 2008. [2](#)
- [15] Saunders Mac Lane and Saunders MacLane. *Categories for the working mathematician*, volume 2. Springer, 1971. [3](#)
- [16] Youssef Mroueh, Stephen Voinea, and Tomaso Poggio. Learning with Group Invariant Features: A Kernel Perspective. *Advances in Neural Information Processing Systems*, 28, 2015. [2](#)
- [17] Walter Rudin. Real and complex analysis. 1974. [2](#)

- [18] Bernhard Schölkopf and Alexander J Smola. *Learning with kernels: support vector machines, regularization, optimization, and beyond*. the MIT Press, 2001. [1](#), [2](#)
- [19] Christian Walder and Olivier Chapelle. Learning with Transformation Invariant Kernels. *Advances in Neural Information Processing Systems*, 20, 2007. [2](#)
- [20] Christopher KI Williams and Carl Edward Rasmussen. *Gaussian processes for Machine Learning*, volume 2. MIT press Cambridge, MA, 2006. [2](#)
- [21] Matthew O. Williams, Clarence W. Rowley, and Ioannis G. Kevrekidis. A kernel-based method for data-driven Koopman spectral analysis, 2015. [2](#)

A Proofs

A.1 Proof for Proposition 2.1

Proof. Fix $x \in X$ and define $k_{G,x} = f_{G^{-1}\Psi(x)}$. Then, for every $f_c \in \mathcal{V}$,

$$\begin{aligned}\langle f_c, k_{G,x} \rangle_G &= \Psi(x)^* G^{-1} G c \\ &= \Psi(x)^* c \\ &= f_c(x).\end{aligned}$$

Thus $k_{G,x}$ represents point evaluation at x , and hence $(\mathcal{V}, \langle \cdot, \cdot \rangle_G)$ is an RKHS satisfying $k_G(x, y) = \Psi(x)^* G^{-1} \Psi(y)$. $\square$

A.2 Proof for Theorem 2.2

Proof. For $f_c \in \mathcal{V}$, $\|Tf_c\|_G^2 = \|f_{Mc}\|_G^2 = c^* M^* G M c$. Hence

$$\|Tf_c\|_G^2 = \|f_c\|_G^2 \quad \text{for every } c \in \mathbb{C}^m \iff M^* G M = G.$$

Suppose first that $G \succ 0$ satisfies $M^* G M = G$. Define $U = G^{1/2} M G^{-1/2}$. Then

$$U^* U = G^{-1/2} M^* G M G^{-1/2} = I,$$

so U is unitary. Therefore, M is similar to a unitary matrix.

Conversely, suppose $M = S D S^{-1}$, where $D = \text{diag}(\lambda_1, \dots, \lambda_m)$, $|\lambda_j| = 1$. Set $G = S^{-*} S^{-1}$. Then $G \succ 0$ and

$$\begin{aligned}M^* G M &= S^{-*} D^* D S^{-1} \\ &= S^{-*} S^{-1} \\ &= G.\end{aligned}$$

The spectral characterization follows. $\square$

A.3 Proof for Theorem 2.4

Proof. Write $G = S^{-*} H S^{-1}$. Then $M^* G M = G \implies D^* H D = H$. Writing $D = \text{diag}(\lambda_1, \dots, \lambda_m)$, the (i, j) entry of this equation is $\bar{\lambda}_i \lambda_j H_{ij} = H_{ij}$. Thus $(\bar{\lambda}_i \lambda_j - 1) H_{ij} = 0$.

Since $|\lambda_i| = |\lambda_j| = 1$, $\bar{\lambda}_i \lambda_j = 1 \iff \lambda_i = \lambda_j$. Hence $H_{ij} = 0$ whenever $\lambda_i \neq \lambda_j$. Therefore H is block diagonal with respect to the distinct eigenspaces $H = \text{diag}(H_1, \dots, H_r)$, where each H_j is Hermitian positive definite. The converse follows immediately by substitution. $\square$

A.4 Proof for Corollary 2.5

Proof. By Proposition 2.1, $k_G(x, y) = \Psi(x)^* G^{-1} \Psi(y)$. For $G = S^{-*} \text{diag}(H_1, \dots, H_r) S^{-1}$, we have

$$G^{-1} = S \text{diag}(H_1^{-1}, \dots, H_r^{-1}) S^*.$$

The result follows from Theorem 2.4. $\square$

A.5 Proof for Corollary 2.6

Proof. Theorem 2.4 identifies $\mathcal{G}(M)$ with the product $\prod_{j=1}^r \mathbb{H}_{++}^{m_j}$. The real vector space of $m_j \times m_j$ Hermitian matrices has dimension m_j^2 , and the positive-definite cone is an open subset thereof. Hence $\dim_{\mathbb{R}} \mathcal{G}(M) = \sum_{j=1}^r m_j^2$. In addition to this, the second equality follows from the injective metric-kernel correspondence. $\square$

B Other important results

Corollary B.1 (Simple spectrum). *Suppose that M has m distinct eigenvalues and $\sigma(M) \subset \mathbb{T}$. Then every invariant geometry has the form*

$$G_{\mathbf{h}} = S^{-*} \text{diag}(h_1, \dots, h_m) S^{-1}, \quad h_j > 0,$$

and the corresponding reproducing kernels are

$$k_{\mathbf{h}}(x, y) = \Psi(x)^* S \text{diag}(h_1^{-1}, \dots, h_m^{-1}) S^* \Psi(y).$$

B.1 Uniqueness of the kernel parameterization

Proposition B.2 (Injectivity of the metric–kernel correspondence). *Suppose that $\psi_1, \dots, \psi_m$ are linearly independent. If $k_{G_1}(x, y) = k_{G_2}(x, y)$ for all $x, y \in X$, then $G_1 = G_2$.*

Proof. The equality of the kernels gives $\Psi(x)^* (G_1^{-1} - G_2^{-1}) \Psi(y) = 0$ for all $x, y \in X$. Since $\psi_1, \dots, \psi_m$ are linearly independent, $\text{span}\{\Psi(x) : x \in X\} = \mathbb{C}^m$. It follows that $G_1^{-1} - G_2^{-1} = 0 \implies G_1 = G_2$. $\square$

In particular, if the spectrum is simple, then $\dim_{\mathbb{R}} \mathcal{G}(M) = m$. At the opposite extreme, if $M = \lambda I$, $|\lambda| = 1$, then $\mathcal{G}(M) = \mathbb{H}_{++}^m$, $\dim_{\mathbb{R}} \mathcal{G}(M) = m^2$. Eventually, the combined application of [Proposition B.2](#) and [Theorem 2.4](#) produces the bijective parameterization

$$\mathfrak{K}_{\Psi}(T) \cong \prod_{j=1}^r \mathbb{H}_{++}^{m_j}.$$

B.2 An operator-adapted Gaussian RKHS

The preceding finite-dimensional characterization also induces operator-adapted kernels within standard kernel families. We illustrate this observation for the Gaussian radial basis function kernel.

Let $Q \in \mathbb{C}^{d \times d}$ be unitary and let $S \in \mathbb{C}^{d \times d}$ be invertible. Define $A = SQS^{-1}$. Although A is similar to a unitary matrix, it does not need to be unitary with respect to the Euclidean inner product. In particular, for a non-unitary choice of S , one generally has

$$A^* A \neq I.$$

Define $G = S^{-*} S^{-1}$. Then $G \succ 0$, and

$$\begin{aligned} A^* G A &= (SQS^{-1})^* S^{-*} S^{-1} (SQS^{-1}) \\ &= S^{-*} Q^* S^* S^{-*} S^{-1} S Q S^{-1} \\ &= S^{-*} Q^* Q S^{-1} \\ &= S^{-*} S^{-1} \\ &= G. \end{aligned}$$

Thus A is unitary with respect to the inner product induced by G .

Consider the Gaussian kernel

$$k_G(x, y) = \exp \left(-\frac{(x - y)^* G (x - y)}{2\sigma^2} \right), \quad x, y \in \mathbb{R}^d,$$

where $\sigma > 0$. Since $A^* G A = G$,

$$\begin{aligned} k_G(Ax, Ay) &= \exp \left(-\frac{(A(x - y))^* G A(x - y)}{2\sigma^2} \right) \\ &= \exp \left(-\frac{(x - y)^* A^* G A(x - y)}{2\sigma^2} \right) \\ &= k_G(x, y). \end{aligned}$$

Hence the kernel is invariant under the transformation A .

Let $\mathcal{K}_A$ denote the Koopman operator

$$(\mathcal{K}_A f)(x) = f(Ax).$$

For the kernel sections $k_x = k_G(\cdot, x)$, the adjoint satisfies

$$\mathcal{K}_A^* k_x = k_{Ax}.$$

Therefore

$$\begin{aligned} \langle \mathcal{K}_A^* k_x, \mathcal{K}_A^* k_y \rangle_{\mathcal{H}_{k_G}} &= \langle k_{Ax}, k_{Ay} \rangle_{\mathcal{H}_{k_G}} \\ &= k_G(Ay, Ax) \\ &= k_G(y, x) \\ &= \langle k_x, k_y \rangle_{\mathcal{H}_{k_G}}. \end{aligned}$$

Since the linear span of the kernel sections is dense in $\mathcal{H}_{k_G}$, $\mathcal{K}_A^*$ extends isometrically to the whole space. As A is invertible, the induced composition action is surjective, and the corresponding Koopman operator is unitary on $\mathcal{H}_{k_G}$.

By contrast, the isotropic Gaussian radial basis function kernel [5] $k_I(x, y) = \exp\left(-\frac{\|x - y\|_2^2}{2\sigma^2}\right)$ satisfies

$$k_I(Ax, Ay) = k_I(x, y)$$

for all x, y only when $A^*A = I$. Thus a transformation that is not Euclidean-unitary need not preserve the standard Gaussian RKHS, while the operator-adapted metric G produces a Gaussian RKHS in which the same transformation acts unitarily.

B.3 Normalization

The invariance relation is homogeneous. Indeed, if $G \in \mathcal{G}(M)$, then $cG \in \mathcal{G}(M)$ for every $c > 0$. Moreover,

$$k_{cG}(x, y) = \frac{1}{c} k_G(x, y).$$

Thus positive scalar multiplication changes the normalization of the Hilbert norm while preserving the underlying function space. It is therefore natural to consider invariant geometries modulo the equivalence relation

$$G_1 \sim G_2 \iff G_2 = cG_1 \text{ for some } c > 0.$$

Define

$$\mathfrak{M}_\Psi(T) = \mathcal{G}(M)/\mathbb{R}_{>0}.$$

Equivalently, one may choose a representative of each equivalence class by imposing, for example, $\text{tr}(G) = 1$. The resulting normalized family has real dimension

$$\dim_{\mathbb{R}} \mathfrak{M}_\Psi(T) = \sum_{j=1}^r m_j^2 - 1.$$

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The claims in the abstract and introduction match the theoretical existence and classification results in [Section 2](#) and the numerical verification in [Sub-section 2.3](#). The finite-dimensional scope of the theory is stated explicitly in the paper.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: The principal limitation is discussed in [Section 3](#): the present classification concerns finite-dimensional function spaces. The infinite-dimensional RKHS problem requires an operator-level analogue of the invariant metric equation and is left beyond the present scope.

Guidelines:

- The answer [\[N/A\]](#) means that the paper has no limitation while the answer [\[No\]](#) means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: The assumptions are stated with the corresponding theoretical results in [Section 2](#), and complete proofs of [Proposition 2.1](#), [Theorem 2.2](#) and [Theorem 2.4](#), and [Corollary 2.5](#) and [Corollary 2.6](#) are provided in [Appendix A](#).

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: The synthetic construction, matrix dimensions, sample size, kernel bandwidth, and invariance-error metric are specified in [Sub-section 2.3](#).

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include experiments.
- If the paper includes experiments, a [\[No\]](#) answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in

some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: The paper does not currently provide an open code repository. The numerical experiment is synthetic and uses no external dataset; its mathematical construction and evaluation criterion are described in [Sub-section 2.3](#).

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The numerical verification is a synthetic, non-training experiment. [Sub-section 2.3](#) specifies the representation dimension $d = 20$, sample size $N = 120$, kernel bandwidth $\gamma = 3$, the invariant metric condition, and the kernel-invariance evaluation criterion. No training/test split or optimizer is involved.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [N/A]

Justification: The experiment is a deterministic numerical verification of an exact algebraic invariance relation rather than a statistical performance comparison. Consequently, error bars, confidence intervals, and statistical significance tests are not applicable to the reported claim.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: The numerical verification is a small synthetic linear-algebra experiment requiring negligible computational resources, but detailed hardware and runtime information are not reported.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research conforms to the NeurIPS Code of Ethics. It is a theoretical and synthetic numerical study and does not involve human subjects, personal data, sensitive datasets, or deployed decision-making systems.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: The work is foundational mathematical research on invariant RKHS geometries and does not introduce a deployed system, dataset, or model with a direct societal application. We therefore do not identify specific positive or negative societal impacts arising directly from the present results.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The work does not release high-risk models, sensitive datasets, or other assets requiring safeguards for responsible release.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [N/A]

Justification: The paper does not use external datasets, pretrained models, or other research assets requiring licensing. Existing scholarly work is credited through the references.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: The paper does not introduce or release a new dataset, pretrained model, benchmark, or other research asset.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The research does not involve crowdsourcing or human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The research does not involve human subjects; therefore, IRB or equivalent human-subjects approval is not applicable.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not used as an important, original, or non-standard component of the methodology, theoretical results, or numerical experiment.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.